

Redis Connect

Version 0.10.3

Table of Contents

Introduction.....	1
Key Terms.....	2
Source Distributions	3
Supported Sources	4
Redis Connect Management.....	5
Production Deployment	8
Configuration Overview	17
Cluster Configuration	18
Job Configuration	27

Introduction

Redis Connect is a distributed platform that enables real-time event streaming, transformation, and propagation of [changed-data events](#) from heterogeneous data platforms to Redis Stack, Redis Cloud, and Redis Enterprise.

Redis Connect effectively captures change data events from source databases and writes that data to Redis. This allows you to keep a Redis database in sync with a variety of source databases. You can then use Redis to serve this data to downstream applications at low latencies.

This document outlines key terms, installation instructions, production readiness guidelines, and definitions for the various configuration directives for Redis Connect clusters and jobs.

- [Key Terms](#)
- [Source Distributions](#)
- [Production Deployment](#)
- [Configuration Overview](#)
- [Cluster Configuration](#)
- [Job Configuration](#)

Key Terms

To understand how Redis Connect works, it's important review some key terms:

Source

A database, such as PostgreSQL, whose data will be replicated to Redis. Redis Connect replicates data from a source database to Redis.

Target

A database to write data to. With Redis Connect, the target is usually Redis.

Job

A stream of change-data events replicating from source to target. For example, you can replicate all changes from a given set of PostgreSQL tables to Redis and maintain consistency between these tables and their Redis representations in real time.

Job Types

Redis Connect supports two types of jobs: **initial load jobs** and **stream jobs**.

- **Initial load jobs** create a point-in-time snapshot of the tables to be replicated and then transfer their data to Redis.
- **Stream jobs** (also known as CDC or "change data capture" jobs) replicate changes from the source tables to Redis as those changes occur.

Instance

A single JVM process running Redis Connect. Because Redis Connect is a distributed platform, it may run as one or more coordinated instances.

Cluster

One or more **instances** of Redis Connect running in a coordinated fashion.

Partition

A way of dividing jobs to scale them horizontally. Each job may be divided into one or more partitions. Partitions are divided automatically among Redis Connect instances. If a Redis Connect instances becomes unavailable, all jobs partitions will be migrated to another available instance.

Source Distributions

Redis Connect releases are distributed on Github.

See the [Redis Connect Release History](#) to download the latest distributions of Redis Connect.

The source distribution contains four relevant folders:

lib

JARs for Redis Connect and its dependencies.

extlib

JARs for Redis Connect custom stages and database drivers not included the Redis Connect distribution (e.g., Oracle and DB2). These database drivers must be provided by the Redis Connect user.

config

Working configuration files, sample payloads for configuring jobs, and Grafana dashboard configurations.

bin

Scripts for running Redis Connect on Linux VMs, in container environments, and on Windows.

Supported Sources

Redis Connect can capture change data from several RDBMS and NoSQL databases. Redis Connect can also load data from CSV and JSON files.

The table below show the complete list of supported sources. Each source database name links to a demo with sample configuration and data.

Under the hood, Redis Connect uses [Debezium](#) to access most of its source databases. In production, it's important that your source database is configured as Debezium expects. Be sure to review the Debezium doc links below.

Table 1. Supported source databases

Database name	Job types	Debezium Docs
DB2	Initial load	Debezium docs
Files	Initial load	n/a
Gemfire	Initial load & Stream	n/a
MongoDB	Initial load & Stream	Debezium docs
MySQL	Initial load & Stream	Debezium docs
Oracle	Initial load & Stream	Debezium docs
PostgreSQL	Initial load & Stream	Debezium docs
Splunk	Initial load	n/a
SQLServer	Initial load & Stream	Debezium docs
Vertica	Initial load	n/a

Redis Connect Management

Once Redis Connect is installed and running, you manage Redis Connect using its REST API or command line interface (CLI).

REST API

Each Redis Connect instance can be configured to expose a REST API with a Swagger interface for ease of use.

The table below shows a few of the commonly-used REST endpoints. For the complete REST API documentation, see the [Redis Connect Swagger API Docs](#).

Table 2. Common REST API Endpoints

Endpoint name	Description	Documentation
Create job	<p>Saves job configuration for the provided <code>jobName</code>.</p> <p>Many of the optional job configuration attributes have default values which can be reviewed once the job configuration is saved or found in the documentation.</p> <p>Example Redis key: <code>{connect}:job:config:jobName</code></p>	Create job Swagger endpoint

Endpoint name	Description	Documentation
Start job	<p>Starts a job, including all job partitions.</p> <p>This includes both initial load and stream jobs. For a job start, all job partitions must be stopped or never before started. There is no guarantee on which cluster instance will claim a job partition, and there is no advantage to initiating this operation from a specific cluster member. Before a start is initiated, a job configuration must be created and a validated to confirm enough remaining capacity exists across the cluster for all job partitions (this does not apply to initial load).</p>	Start job Swagger endpoint
Stop job	<p>Stops a job, including all job partitions.</p> <p>You cannot stop initial load jobs since they are removed automatically upon completion. For a job to be stopped, all job partitions must be active. Job claims, metrics, and checkpoints will all be preserved upon a job stop. This ensures that the job can later resume from where it was stopped.</p>	Stop job Swagger endpoint
Job claim status	<p>Returns job claims from across the cluster that match the requested jobStatus.</p> <p>Valid job statuses are staged, stopped, and all.</p> <p>Use this endpoint to see which jobs the cluster is managing, and their status.</p>	Job claim status Swagger endpoint

REST API Security

In a production environment, the Swagger API may require an open port in a firewall. By default, the API is available on port 8282, but this is configurable.

If you are running multiple Redis Connect instances on the same server, each instance will require a different port for its REST API.

CLI

Redis Connect includes a command line interface that exposes the management functions provided by the REST API. You can start a CLI instance from the command line as follows:

```
$ ./bin/redisconnect.sh cli
```

Once the CLI has started, type **help** to see the available commands.

Production Deployment

Redis Connect is deployed as one or more JVM instances coordinated as a cluster. Below are recommendations for running Redis Connect in production.

Environment

Redis Connect can be deployed on physical servers, virtual machines, or using Docker or any Kubernetes-based environment.

The minimum resource requirements per Redis Connect instance are as follows:

- 4 CPU cores
- 2 GB memory
- 20 GB of free disk space
- 1 Gbps networking

We recommend allocating one thread per job partition. You can deploy more than one Redis Connect instance on a single VM. The number of Redis Connect instances that can effectively be deployed on a single machine or VM will depend on that VM's physical memory and number of CPU cores. However, for high availability, you must deploy your Redis Connect instances across more than one physical server or VM. See [High Availability](#) below for more details.

Operating System and JVM

Redis Connect can run on any operating system hosting a Java runtime environment. However, for production deployments, we recommend Linux.

Redis Connect is supported on Java versions 11 and greater.

For information on deploying to Kubernetes environments, see the [Redis Connect Kubernetes documentation](#).

Environment Variables

Redis Connect recognizes and depends upon several environment variables. You can see example of these in the startup scripts included in the Redis Connect distribution.

- `REDISCONNECT_MIN_JAVA_VERSION="11"`
- `REDISCONNECT_HOME="${REDIS_CONNECT_HOME_DIR}"`
- `REDISCONNECT_JOB_MANAGER_CONFIG_PATH="$REDISCONNECT_HOME/config/jobmanager.properties"`
- `REDISCONNECT_LOGBACK_CONFIG="$REDISCONNECT_HOME/config/logback.xml"`
- `REDISCONNECT_LOGBACK_CLI_CONFIG="$REDISCONNECT_HOME/config/logback-cli.xml"`
- `REDISCONNECT_JAVA_OPTIONS="-XX:+HeapDumpOnOutOfMemoryError -Xms1g -Xmx2g"`
- `REDISCONNECT_EXTLIB_DIR="$REDISCONNECT_HOME/extlib"`

- `REDISCONNECT_LIB_DIR="$REDISCONNECT_HOME/lib/*:$REDISCONNECT_EXTLIB_DIR/*"`

JVM Flags

We recommend the following JVM flags for each Redis Connect JVM instance:

- Enable heap dump on OOM: `-XX:+HeapDumpOnOutOfMemoryError`
- Min heap size of 1 GB: `-Xms1g`
- Max heap size of 2 GB: `-Xmx2g`

To set Redis Connect's JVM options, use the `REDISCONNECT_JAVA_OPTIONS` environment variable:

```
REDISCONNECT_JAVA_OPTIONS="-XX:+HeapDumpOnOutOfMemoryError -Xms1g -Xmx2g"
```

Redis Requirements

Redis Connect requires a working [Redis Enterprise Software](#) or [Redis Cloud](#) installation.

We recommend provisioning two Redis databases. These database can reside in the same Redis Enterprise cluster.

The first database will serve as your Redis Connect cluster's configuration store. This database should be configured as follows:

- Data persistence enabled (RDF + AOF every second)
- [High availability](#) enabled
- ACLs enabled (See [Security](#) below)

The second database will serve as the target for replication from the source database. Since this is an operational database receiving change-data events, the sizing for this database depends on the sizes of the tables being replicated and on the volume of change-data events. In all cases, we still recommend:

- [High availability](#) enabled
- ACLs enabled (See [Security](#) below)

Logging

Redis Connect uses [Logback](#) for logging. See your Redis Connect distribution's `config/logback.xml` for a sample Logback configuration file.

Redis connect locates its logging config files using the following environment variables:

- `REDISCONNECT_LOGBACK_CONFIG="$REDISCONNECT_HOME/config/logback.xml"`
- `REDISCONNECT_LOGBACK_CLI_CONFIG="$REDISCONNECT_HOME/config/logback-cli.xml"`

Redis Connect has been designed to provide descriptive logs to make troubleshooting easier.

If you need to change your application log level at runtime, you can do this using the [REST API](#). See the [loglevel REST endpoint documentation](#) for details.

Note that log level changes are not global; they apply only to the instance whose REST API you are connected to. To change the log level for a given instance, connect directly to that instance's REST API.

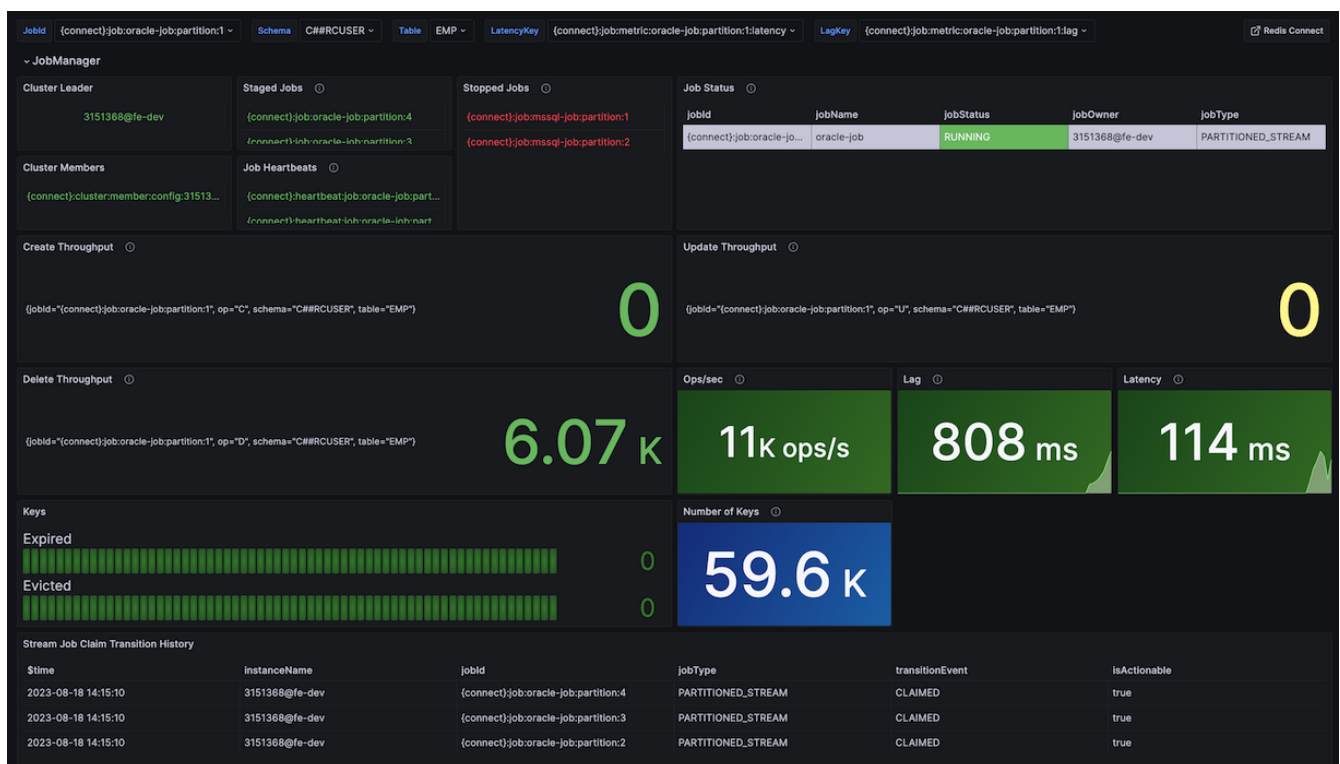
Monitoring

Redis Connect publishes performance metrics to job manager's configuration database. This takes advantage of [Redis' time series capabilities](#).

You can view these metrics in Grafana using the [Redis Datasource for Grafana](#).

Your Redis Connect distribution includes a pre-configured dashboard for viewing key operational metrics. You can find this dashboard and its Grafana configuration at [config/samples/dashboard](#), and this configuration can be modified to fit your monitoring requirements.

The sample dashboard looks like this:



Important metrics include the following:

Staged jobs

The list of all configured jobs in the Redis Connect cluster.

Job heartbeats

Heartbeats for the currently active job partitions. Each active partition records a heartbeat every second, indicating that the job is actively being run by a Redis Connect instance.

Ops per second

The number of write operations against the target database completed per second.

Lag

The average elapsed time between the moment a change event is published to the source database and the moment that event is written to the target database (i.e., Redis). High lag values may indicate that your Redis Connect cluster is failing to keep up with the volume of CDC changes. Note that this metric only applies to stream jobs.

Latency

The average amount of time, in milliseconds, that it takes to publish a change event from the source database to Redis.

Security

Redis Connect should be run in a secure network environment. Because Redis Connect connects to source and target databases, those connections **may** be secured with TLS/SSL and **must** be secured using authentication credentials.

To support additional security requirements, Redis Connect supports file-based credential rotation, optionally powered by secrets management frameworks like Hashicorp Vault.

TLS/SSL Support

Most database systems support one-way and mutual TLS, and TLS authentication is often required in production environments.

Redis Connect manages certificate-based authentication using Java's KeyStore and TrustStore support. To configure the Java KeyStore and TrustStore, see the [KeyStore and TrustStore configuration reference](#).

To configure Redis Enterprise with TLS, see the documentation:

- [TLS with Redis Cloud](#)
- [TLS with Redis Enterprise Software](#)

Database ACLs

The Redis job manager and the job target databases should have ACLs enabled. To keep your Redis databases secure and prevent accidental misconfigurations, we recommend the following policies:

1. Create separate Redis database users for your Redis Connect instances, Redis Connect administrators, and developers.
2. Developers should be provided with read-only access to these Redis databases.
3. Ensure that any application user connecting Redis Connect's Redis databases has **dangerous** commands disabled.
4. Users should not have permission to delete keys starting with "{connect}". This prevents

accidental deletion of important configuration.

Authentication Credentials

Redis Connect gets its database authentication credentials from properties files stored on the filesystem.

The file `redisconnect_credentials_jobmanager.properties` is common to all Redis Connect instances. This file provides the following authentication credentials:

- Username and password for the Redis database used to store Redis Connect's configuration
- Password for the certificate TrustStore (when applicable)
- Password for the certificate KeyStore (when applicable)
- Password for an SMTP mail server, when email alerts are enabled

In addition to the job manager's credentials, Redis Connect requires two credential files for each job: one for the source database and another for the target Redis database.

Source database credential files are named according to this scheme:

```
redisconnect_credentials_[SOURCE_DB_NAME]_[JOB_NAME].properties
```

For example, if you have a job where PostgreSQL is the source database, and the job is called "user-replication", then you will need a properties file called:

```
redisconnect_credentials_PostgreSQL_user-replication.properties
```

The **target database credentials files** are named as follows:

```
redisconnect_credentials_redis_[JOB_NAME].properties
```

For job called "user-replication", the Redis target database credentials file will be named:

```
redisconnect_credentials_redis_user-replication.properties
```

You can see examples of credentials files in the `config/samples/credentials` directory of the Redis Connect distribution.

Securing Credentials Files

Because credentials files store sensitive information in plain text on the filesystem, these files must have strict permissions. Only authorized users should be able to read and write to these files.

As with private key files, these credentials files should be user read/write only (e.g., `chmod 600`).

Credential Rotation

Redis Connect instances listen for changes to the credentials files. When the contents of these files changes, Redis Connect will read the changes to ensure that database connectivity is maintained.

To enable support for credential rotation in Redis Connection, open `jobmanager.properties` and ensure that the credential rotation directives are uncommented. The following configuration enabled the credential file listener and sets it to check for changes every 60 seconds:

- `credentials.rotation.event.listener.enabled=true`
- `credentials.rotation.event.listener.interval=60000`

Secrets Management

You can use a secrets management framework such as Hashicorp Vault to securely store and rotate credentials. To use a secrets management framework:

1. Ensure that framework can write property files in a given mounted path on your filesystem.
2. In `jobmanager.properties`, set the `credentials.dir.path` to this mounted path.

High availability

For high-availability, we recommend that you employ $n+1$ redundancy as part of your Redis Connect cluster. In this case, n is defined as the minimum number of Redis Connect instances required to provision all of the jobs that you intend to run.

We also recommend that you deploy each Redis Connect instance on a separate VM. When running in the cloud, you should ensure that your Redis Connect instances are distributed across cloud availability zones.

Instance failure behavior

In the event of a Redis Connect instance (JVM) failure, the heartbeat lease of each job partition owned by that instance expires. Once expired, the job reaper will identify each partition without a heartbeat and immediately publish a job claim request.

Even if a JVM was restarted on the same machine, on which it just failed, it would still need to compete with the other instances with available capacity to claim the job partitions.

It's important to understand each job's capacity requirements and related settings to calculate how many instances are needed. For example, if a job with 4 partitions and a setting of `maxPartitionsPerClusterMember=2` is deployed across a 2-instance cluster, then even if there are 20 available capacity on each Redis Connect instance, they would be blocked from claiming more of this job's unclaimed partitions in the event of a node failure.

It's also important to set the Job Manager property `job.claim.max.capacity` appropriate to the desired redundancy requirement with respect to all job partitions planned for deployment across the cluster. For example, if there are 3 cluster nodes each with a goal to deploy a job with 4 partitions, then in order to achieve $n+1$ redundancy, each node would have to configure

`job.claim.max.capacity=2`, or more, since the loss of one node would still allow for all 4 partitions to be claimed.

Instance failures

For high-availability, it is recommended to deploy each Redis Connect instance on a separate VM or cloud availability zone. For on-premises deployments, extra care should be taken to ensure that each VM is deployed on different underlying server racks to avoid a single point of failure.

Each VM / cloud availability zone should have access to each job's source, job's target, and the Redis Connect Job Manager database.

For deployment on Kubernetes, see the [Redis Connect Kubernetes Docs](#).

Network partitions

If a network partition occurs, then one side or the other (or both) of the partition needs to stop responding to requests to maintain the consistency guarantee. If both sides continue to respond to reads and writes while they are unable to communicate with each other, they will diverge and no longer be consistent. This state where both sides of the partition remain available is called "split brain".

To avoid "split brain", Redis Connect jobs make a call to the Job Management database as the initial stage of their producer's source event-loop to check if their claim is still valid. Here's why that matters. A job cannot be claimed unless the Job Reaper identifies a staged job without a heartbeat, so until that lease expires, the claim is intact. Once the heartbeat expires due to the network partition, the Job Reaper will first remove the claim from the Job Management database before publishing a new job claim request. So even if the original Redis Connect instance has not died, it cannot continue to process additional batches once its claim is removed nor can it regain its ownership by default when the network partition is remedied.

By relying on Redis Enterprise as its source of truth, this cluster architecture is guaranteed to always have a single job owner at any one point-in-time. This is because the metadata stored within Redis Enterprise has its own split-brain protection.

Source Connection Failures

Connections fail for many reasons both intermittently and for prolonged outages. Since Redis Connect does not know the cause of a connection issue, the cluster will always assume that the connection failure is intermittent and attempt to re-establish the connection.

For this purpose, it's important to set the following configurations within the job source:

- `sourceConnectionMaxRetryAttempts`
- `sourceConnectionMaxRetryDuration`
- `sourceConnectionRetryDelayInterval`
- `sourceConnectionRetryMaxDelayInterval`
- `sourceConnectionRetryDelayFactor`

Here's an example of their use:

- `sourceConnectionMaxRetryAttempts=3`
- `sourceConnectionMaxRetryDuration=10`
- `sourceConnectionRetryDelayInterval=60`
- `sourceConnectionRetryMaxDelayInterval=600`
- `sourceConnectionRetryDelayFactor=3`

Based on the above, if the first attempt to connect fails, then the process waits 60 seconds before attempting its first retry.

If the retry fails, then the process waits $60 * 3 = 180$ seconds for its second retry, and so on, up until the max delay interval between retries.

Target Redis Connection Failures

Just like sources, Redis target connections have a retry mechanism. However, these retries are built into the underlying Redis client (a.k.a., [Lettuce](#)). For production environments, we recommend using the default Redis connection settings.

Target Redis Slow Consumer

It's possible that the rate of incoming changed-data events and outgoing sink commits becomes unbalanced due to a slowdown on the target database, usually because of network congestion. To avoid overwhelming the job pipeline queue, the combination of the following settings helps to provide back-pressure support:

- `slowConsumerMaxRetryAttempts`
- `intermittentEventSleepDuration`

Once the pipeline queue becomes full, a circuit breaker will be triggered to disable the producer's polling event-loop until the requisite pipeline capacity becomes available. While the pipeline is iterating through its queued events, the job's producer will periodically attempt to publish the remainder of its current batch.

In between each attempt, the producer will sleep (based on `intermittentEventSleepDuration`), which allows the pipeline time to catch up. This loop will continue until either the required capacity becomes available or the `slowConsumerMaxRetryAttempts` limit is reached. Once reached, the job will stop.

In production, we recommend that you set `slowConsumerMaxRetryAttempts` to a reasonable setting instead of making it unlimited or disabling it. This is where proper capacity estimation and partitioning will become critical since you don't want to stop the job unnecessarily due to a network disruption nor large spike in volume. In other words, even if a single partition can handle the entire throughput on a regular basis, it might be worthwhile to partition the deployment so that a sudden spike in throughput or a slow consumer does not bring down the job.

Recovery Time Considerations

Redis Connect's job management processes are scheduled threads that periodically wake up to perform their service. While it's not critical to understand the internals of the cluster architecture, it is relevant to be aware of the balance between job claim heartbeats and the Job Reaper service.

As mentioned previously, when a job partition is claimed it begins a heartbeat lease which has a TTL set by `job.claim.heartbeat.lease.renewal.ttl`. If the lease cannot be renewed before expiry, the Job Reaper service will remove the current owner's claim based on the assumption that its Redis Connect instance is no longer alive.

This service is performed based on a fixed interval set by `job.reap.attempt.interval`. Once the job claim is removed, each eligible Redis Connect instance will attempt to claim the job partition. This Job Claim service is also based on a fixed interval set by `job.claim.attempt.interval`.

Since each interval is started at a different point-in-time (and some services have more work to do than others), it's impossible to synchronize all the moving parts. However, it is critical to avoid scenarios that can cause unexpected behavior by giving enough buffer between associated services. For example, `job.reap.attempt.interval` must be at least 2 seconds more than `job.claim.attempt.interval` to avoid sending multiple job claim requests for the same partition. This relates to recovery because each default interval can potentially add up to over a minute in recovery time inclusive of the actual start job process. Therefore, if instant recovery is critical to production SLAs, our recommendation is to consider the balance between separate services and as well as the stability of your network. If these settings are set too low, it's possible that a network disruption can unnecessarily start the recovery process on a healthy job.

See [Cluster Configuration](#) for details on setting these properties.

Configuration Overview

Redis Connect's main configuration file is called "jobmanager.properties". You can find an example of this file in your distribution's "/config" directory.

The location of this file is provided to Redis Connect with the `REDISCONNECT_JOB_MANAGER_CONFIG_PATH` environment variable. Here's an example of its use:

- `REDISCONNECT_JOB_MANAGER_CONFIG_PATH="$REDISCONNECT_HOME/config/jobmanager.properties"`

The most basic Redis Connect configuration includes the url of the Redis database used to store this state. This URL is stored at the parameter `redis.connection.url` in `jobmanager.properties`.

Job Configuration

Redis Connect runs one or more **change-data capture jobs**. Each job is defined by a JSON configuration payload submitted using the REST API or CLI. To see examples of this JSON configuration, open the `config/samples/payloads` folder in the Redis Connect distribution.

To save a job configuration, construct your JSON payload and then submit it using the REST API's [create job endpoint](#). When you submit this configuration, you will also provide a unique job name. This job name will be used in the naming of your credentials files.

Credentials files

Credentials files store authentication credentials for the databases used by Redis Connect.

The location of these credentials files is specified at the `credentials.dir.path` properties in `jobmanager.properties`. To learn about which credentials are required, see the [authentication credentials](#) subsection.

Cluster Configuration

This section describes the configuration directives in `jobmanager.properties`. Each of the subsections below corresponds to a configuration block in `jobmanager.properties`. To see a sample `jobmanager.properties` file, open `config/jobmanager.properties` in your Redis Connect distribution.

Cluster properties

Table 3. Cluster properties

Property name	Type	Description	Default
<code>cluster.name</code>	String	Metadata purposes only. Non-functional	default
<code>cluster.leader.heartbeat.lease.renewal.ttl</code>	Integer	TTL (Time-to-Live) which is renewed upon each <code>cluster.election.attempt.interval</code> iteration by the cluster leader. Measured in milliseconds with a minimum of 1 second (1000 ms).	5000
<code>cluster.election.attempt.interval</code>	Integer	Fixed rate scheduled thread which either renews or elects a new cluster leader. Runs on each Redis Connect Instance (JVM) when <code>job.manager.services.enabled=true</code> . Measured in milliseconds with a minimum of 1 second (1000 ms)	5000
<code>cluster.timeseries.metrics.enabled</code>	Boolean	Enables creation of a scheduled thread for job metrics reporting to RedisTimeSeries within the Job Management database.	false

Job manager services properties

Table 4. Job manager services properties

Property name	Type	Description	Default
job.manager.services.enabled	Boolean	<p>Enables creation of scheduled thread(s) to participate in cluster leader elections, facilitate REST API / CLI (Job Manager service), and identify staged jobs without a heartbeat lease (Job Reaper service).</p> <p>When this property is disabled, the Redis Connect instance may still participate in job execution and job claim attempts (Job Claimer service).</p>	true
job.manager.services.threadpool.size	Integer	For non-production deployments, one thread is adequate. In production, we recommend two threads.	2
job.reap.attempt.interval	Integer	<p>The interval between attempts to identify staged jobs without a heartbeat lease. Implemented as a scheduled thread that runs on each Redis Connect Instance (JVM) when <code>job.manager.services.enabled=true</code>.</p> <p>Measured in milliseconds with a minimum of 1 second (1000 ms)</p>	7000

Property name	Type	Description	Default
job.claim.service.enabled		<p>Enables creation of scheduled thread(s) to attempt to claim ownership for UNASSIGNED staged jobs (Job Claimer Service), job execution, and job-level metrics reporting (Metrics Reporter service).</p> <p>When this property is disabled, the Redis Connect instance may still participate in cluster leader election, facilitate REST API / CLI, and perform Job Reaper services.</p>	true
job.claim.attempt.interval	Integer	<p>Interval at which this scheduled thread attempts to claim ownership for UNASSIGNED staged jobs.</p> <p>Runs on each Redis Connect Instance (JVM) when <code>job.claim.service.enabled=true</code>.</p> <p>Measured in milliseconds with a minimum of 1 second (1000 ms)</p>	5000
job.claim.batch.size.per.attempt	Long	Specifies how many jobs can be claimed per attempt interval. If a sparse topology across many Redis Connect instances is desired, then lowering this interval is recommended.	4

Property name	Type	Description	Default
job.claim.max.capacity	Integer	Specifies the maximum number of jobs that a single Redis Connect instance can claim at any given time.	4
job.claim.heartbeat.lease.renewal.ttl	Integer	<p>TTL (Time-to-Live) which is renewed upon each iteration of a fixed rate scheduled thread that shares its value.</p> <p>Measured in milliseconds with a minimum of 1 second (1000 ms)</p>	10000

REST API Properties

Table 5. REST API properties

Property name	Type	Description	Default
rest.api.enabled	Boolean	Instantiates an embedded Spring Boot Application to host the REST API and CLI.	true
rest.api.port	Integer	<p>Specifies the port used for the REST API (and SWAGGER) powered by an embedded Spring Boot Application.</p> <p>If you are running multiple Redis Connect instances on the same server, each instance will require a different port for its REST API.</p>	8282

Job Management Database Properties

Table 6. Job Management Database Properties

Property name	Type	Description	Default
redis.connection.url	String	<p>A Redis URI indicating which Redis server to use for job management.</p> <p>For the Redis URI spec, the Lettuce documentation.</p>	n/a
redis.connection.insecure	Boolean	<p>Passed to Lettuce's <code>RedisURI.verifyPeer</code>.</p> <p>If true then <code>verifyMode=FULL</code>. Otherwise, if false, then <code>verifyMode=NONE</code>.</p> <p>When peer verification is disabled, Lettuce uses Netty's <code>InsecureTrustManagerFactory.INSTANCE</code> as the trust manager factory. Its javadoc notes that it should never be used in production and that it is purely for testing purposes.</p>	false
redis.connection.timeout.duration	Integer	The timeout is canceled upon command completion/cancellation. Measured in seconds.	1
redis.connection.auto.reconnect	Boolean	<p>Determine whether the driver will attempt to automatically reconnect to Redis.</p> <p>When enabled, then on disconnect, the client will try to reconnect, activate the connection and re-issue any queued commands.</p>	true

Property name	Type	Description	Default
redis.connection.suspended.reconnect.on.protocol.failure	Boolean	<p>When set to true, reconnect will be suspended on protocol errors.</p> <p>The reconnect itself has two phases: Socket connection and protocol/connection activation. In case a connection timeout occurs, a connection reset, or host lookup fails, this does not affect the cancellation of commands. In contrast, where the protocol/connection activation fails due to SSL errors or PING before activating connection failure, queued commands are canceled.</p>	true
redis.connection.sslEnabled	Boolean	Enables SSL for one-way or mutual authentication. If this flag is set to false , TrustStore and KeyStore will not be passed to the client.	false
truststore.file.path	String	File path of the Java TrustStore (containing certificates trusted by the client)	n/a
keystore.file.path	String	File path of the Java KeyStore, which stores private key entries, certificates with public keys, or any other secret keys used for various cryptographic purposes.	n/a

Property name	Type	Description	Default
credentials.dir.path	String	<p>The name of the directory containing the Redis Connect credentials file. This directory path must include a properties file named <code>redisconnect_credentials_jobmanager.properties</code>.</p> <p>Redis Connect never caches or persists credentials. Therefore, on each connection with the source, target, or job manager database, the credentials are read from a file. This enhances security and allows for seamless credential rotations and integration with secret management frameworks such as HashiCorp Vault.</p>	../config/ samples/ credentials
credentials.rotation.event.listener.enabled	Boolean	<p>When set to <code>true</code>, a listener will be created on the <code>redisconnect_credentials_jobmanager.properties</code> file within the <code>credentials.dir.path</code> to rotate credentials when they change.</p> <p>This lets you rotate credentials without restarting your Redis Connect instance.</p>	false

Property name	Type	Description	Default
credentials.rotation.event.listener.interval	Integer	<p>When <code>credentials.rotation.event.listener.enabled</code> is set to <code>true</code>, this flag sets the frequency at which is scanned for changes.</p> <p>Measured in milliseconds with a minimum of 60 seconds (60000 ms)</p>	60000

Email Alerting Properties

Table 7. Email Alerting Properties

Property name	Type	Description	Default
mail.alert.enabled	Boolean	Enables email alerts when any error forces a job to stop.	false
mail.smtp.host	String	Hostname of the outgoing mail server.	smtp.gmail.com
mail.smtp.port	Integer	Set the non-SSL port number of the outgoing mail server.	587
mail.smtp.start.tls.enable	Boolean	<p>Set or disable STARTTLS encryption.</p> <p>StartTLS is an extension of the SMTP protocol that tells the email server that the email client wants to use a secure connection using TLS or SSL.</p>	true
mail.smtp.start.tls.required	Boolean	Set or disable the required STARTTLS encryption.	false

Property name	Type	Description	Default
mail.to	String	<p>The email address to send alerts to.</p> <p>This email address will also be used as the personal name.</p> <p>Multiple recipients can be added by delimiting them with a comma.</p>	n/a
mail.debug	Boolean	Set session debugging on or off.	false

Job Configuration

This section describes the fields in job configuration payloads.

Each of the subsections below corresponds to a configuration block in `jobmanager.properties`. To see a sample `jobmanager.properties` file, open `config/jobmanager.properties` in your Redis Connect distribution.

Job properties

Table 8. Job properties

Property name	Type	Description	Constraints	Default
jobName	String	<p>Unique name which is used to derive all other Redis metadata keys related to the job execution workflow.</p> <p>jobName should not be confused with jobId. jobIds are created as part of a job claim. They add-on a namespace to the jobName to identify the jobType and partitionId (if jobType=PARTITIONED_STREAM).</p> <p>When jobName is used in logging or administrative processes (i.e., stopJob), the jobName represents ALL job partitions.</p>	<p>min=4, max=50</p> <p>regex pattern="^[\\s<>(){}\\[\\]\\\"'/\\\\\\\\;`!@#\$\$%&*]*\$"</p>	n/a

Property name	Type	Description	Constraints	Default
partitions	Integer	<p>Indicates how many partitions to create during startJob process. This attribute is ONLY used to partition a job with jobType=PARTITIONED_STREAM. Not jobType=LOAD.</p> <p>CAUTION: Once a job has started, and job claims are created, a job cannot be repartitioned without deleting all job claims and existing checkpoints. Please reach out to Support to assist with the migration of checkpoints to avoid undesired outcomes.</p>	min=1	1

Property name	Type	Description	Constraints	Default
maxPartitionsPerClusterMember	Integer	<p>The number of job partitions that can be claimed, and executed, on the same Redis Connect instance (JVM).</p> <p>If the limit forces partitions to span more instances than are currently deployed, then the job will not be able to start nor migrate.</p> <p>For example, if maxPartitionsPerClusterMember=1 and partitions=3, then the Redis Connect cluster will require at least 3 instances (JVMs) each with at least 1 available capacity to claim a job partition.</p> <p>This is not a global limit; it is only specific at the job level. 0 represents no limit.</p>	min=0	0
pipeline	JSON Object	See Job Pipeline Properties	Not Null	n/a
source	JSON Object	See Job Source Properties	Not Null	n/a

Job Pipeline Properties

Table 9. Job Pipeline Properties

Property name	Type	Description	Constraints	Default
pipelineBuffer Size	Integer	<p>Redis Connect's pipeline is powered by the LMAX Disruptor library (High Performance Inter-Thread Messaging).</p> <p>The buffer size sets the number of slots allocated within the Disruptor's internal ring buffer "queue".</p> <p>Increasing the buffer size will impact the JVM heap space required to store all transient changed data events within the queue. For most cases, this can be left as default.</p>	min=1024 Must be a power of 2	4096
preprocessorName	String	<p>Functional interface (Consumer) that can be run before changed-data events are transformed and published to the pipeline.</p> <p>This is currently not extendable by end users.</p>		n/a

Property name	Type	Description	Constraints	Default
postprocessorName	String	Functional interface (Consumer) that can be run after changed-data events are transformed and published to the pipeline. This is currently not extendable by end users.		n/a
stages	Job Pipeline Stage[]	See Job Pipeline Stage Properties		

Job Pipeline Stage Properties

Table 10. Job Pipeline Stage Properties

Property name	Type	Description	Constraints	Default
stageName	String	Unique name which is used as an exact match reference to a custom-built target sink or a user-defined custom stage.		n/a
index	Integer	Specifies the sequence in which the stages of the pipeline should be orchestrated.	min=1 Begins with 1 and each subsequent index should increment by 1	n/a

Property name	Type	Description	Constraints	Default
metricsEnabled	Boolean	When enabled, the target sink stage will report throughput and latency related metrics for persistence in RedisTimeSeries. This can subsequently be visualized in Grafana.		false
metricsRetentionInHours	Long	Maximum duration for metrics samples as compared to the highest reported timestamp before they expire. Measured in hours.	min=1	4
checkpointStageIndicator	Boolean	Indicates which sink will be responsible for committing the checkpoint to the target database. This is typically performed by the last stage of the pipeline and, often times, it is the only stage in the pipeline.	Job pipeline can only have a single stage with <code>checkpointStageIndicator=true</code>	false

Property name	Type	Description	Constraints	Default
CheckpointTransactionsEnabled	Boolean	<p>Although the producer's polling event loop enqueues changed-data events in batches, each event is processed individually through the pipeline. This is because Redis Connect updates the checkpoint at the changed-data event level and not the batch.</p> <p>When enabled, the checkpoint will be committed as part of an atomic Redis transaction. This eliminates consistency issues and improves performance. Rollback capability is built in to handle any failure scenarios during the transaction so that no data will be lost.</p> <p>When disabled, the checkpoint will be committed after the the changed-data events are written. This adds another network round trip for each changed-data event.</p>	Distributed checkpoints require RedisSearch. We use RedisSearch to index checkpoint keys so that recovery from the latest checkpoint is immediate.	false

Property name	Type	Description	Constraints	Default
keyPrefix	String	Adds a prefix to the target Redis key before the tableName and composition of targetKey enabled columns.		
userDefinedType		<p>To create a custom stage, a factory interface must be extended so that Redis Connect can have visibility to it from a class loading perspective. See section X.X.X.</p> <p>The interface will force the user to create a <code>getType()</code> method which returns a unique String to represent the custom factory. This property must exactly match that custom unique String so that Redis Connect can properly discover and handle it as a custom stage.</p>		
database		See Database Properties (Section 4.4) Configuration for all target database configuration.		

Property name	Type	Description	Constraints	Default
checkpointDatabase		See Database Properties Checkpoint database configuration. This is only required if Redis is not the target destination, which is only supported for Splunk.		

Job Source Properties

Table 11. Job Source Properties

Property name	Type	Description	Constraints	Default
pollSourceInterval	Long	Fixed rate interval representing how long to pause the producer's polling event loop if no new change events were found in the batch. Measured in milliseconds.	min=5	50
batchSize	Integer	Maximum # of events to dequeue from the source-event-queue AND maximum # of events to query from the source transaction log/table/queue upon each interval of the producer's polling event loop.	min=1	500

Property name	Type	Description	Constraints	Default
sourceTransactionTimeSequenceEnabled	Boolean	When enabled, the source commit/transaction timestamp (and sequence# if the timestamp is the same) will be used to calculate latency metrics and passed along as metadata for Redis Streams sink(s).	n/a	false
slowConsumerMaxRetryAttempts	Integer	<p>-1 = UNLIMITED</p> <p>0 = DISABLED</p> <p>1+ = MAX_ATTEMPTS</p> <p>Used as part of back-pressure support for the data pipeline in the event of a slow consumer. If the maximum attempts limit is reached, the job will be stopped for purposes of manual intervention.</p>	min=-1	50

Property name	Type	Description	Constraints	Default
intermittentEventSleepDuration	Integer	Used as part of back-pressure support for the data pipeline in the event of a slow consumer or the circuit breaker is open. Forces the event loop to pause for the configured duration of time. Measured in milliseconds.	min=0	3
sourceConnectionMaxRetryAttempts	Integer	0 = DISABLED 1+ = MAX_ATTEMPTS Maximum retry attempts to reconnect with the source in the event that a connection is lost.	min=0	3
sourceConnectionMaxRetryDuration	Integer	In addition to sourceConnectionMaxRetryAttempts, you can also add a max duration, after which retries will stop if the max attempts haven't already been reached. Measured in minutes.	min=1	5

Property name	Type	Description	Constraints	Default
sourceConnectionRetryDelayInterval	Long	Fixed delay in between sourceConnectionMaxRetryAttempts. Measured in seconds.	min=0 sourceConnectionRetryDelayInterval must be < than sourceConnectionRetryDelayInterval sourceConnectionRetryDelayInterval must be < than sourceConnectionMaxRetryDuration	60
sourceConnectionRetryMaxDelayInterval	Long	Provides an upper bound to calculate the delay interval when sourceConnectionRetryDelayFactor is enabled. Measured in seconds.	min=0	240
sourceConnectionRetryDelayFactor	Integer	0 = DISABLED 1+ = DELAY_FACTOR Factor by which delays are exponentially increased after each source connection retry attempt.	min=0	2
database	Object	See Database Properties Configuration for all source databases.	Not Null	n/a

Property name	Type	Description	Constraints	Default
tables	Map<String, Table>	<p>See Job Source Table Column Properties</p> <p>Configuration for all source tables/collections/regions/logs properties.</p> <p>Each table within the map requires a unique name which will be used as part of target key composition.</p>	Not Null	n/a

Job Source Database Properties

See [Database Properties](#).

Job Source Table Properties

Table 12. Job Source Table Properties

Property name	Type	Description	Constraints	Default
autoConfigColumnsEnabled	Boolean	<p>When enabled, source metadata is queried during the (re)start process to determine sourceColumn names so users do not need to enumerate each within the column's configuration.</p> <p>The columns configuration can be used to override source metadata (i.e., targetName, type, etc.). However, targetKey designation cannot be overridden since only the source table's primary key will be used.</p> <p>This is a common configuration in POCs and development environments since the design of Redis key names are less important than in production. It also allows for less knowledge about the source table schema.</p> <p>This is only supported for RDB sources.</p>	n/a	false

Property name	Type	Description	Constraints	Default
dynamicSchemaEnabled	Boolean	When enabled, columns that are not provided in the columns configuration will be passed through, as-is, to the target. This is currently only supported for MongoDB, Redis Streams Broker, and Files.	n/a	false
prefixTableNameToTargetKeyEnabled	Boolean	When enabled, adds the tableName (defined in the tables configuration) as a prefix to the target Redis key before all other targetKey enabled columns are computed and applied.	n/a	false

Property name	Type	Description	Constraints	Default
deleteOnPrimaryKeyUpdate Enabled	Boolean	<p>When enabled, if the primary key is changed at the source, then an additional operation to DELETE the existing target key will accompany the UPDATE event.</p> <p>This is only supported for RDB sources since primary key changes require a delete and insert of a new row.</p> <p>The DELETE event shares an offset with the UPDATE event both at the source and checkpoint. Redis Connect will handle them within a single pipeline iteration.</p>	n/a	true

Property name	Type	Description	Constraints	Default
changedColumnsOnlyEnabled	Boolean	<p>When enabled, only allows changed (delta) column values to be replicated to the target. This does not include targetKey column(s) which cannot be bypassed. When disabled, all column values will be replicated to the target unless they are individually bypassed at the column-level using changedColumnOnlyEnabled. (See Section 4.2.2)</p> <p>When enabled, the column-level changedColumnOnlyEnabled flag will be overridden for all columns other than those designated as targetKey(s).</p> <p>This is currently only supported for RDB sources.</p>	n/a	false
columns	Job Source Table Column[]	See Job Source Table Column Properties .	n/a	Null
initialLoad	Initial Load	See Job Source Table Initial Load Properties .	n/a	Null

Job Source Table Column Properties

Table 13. Job Source Table Column Properties

Property name	Type	Description	Constraints	Default
targetKey	Boolean	Designates this column's value as part of the target's key composition process. When more than one column is designated, the order in which they are listed will impact the order in which they are appended to the key.	n/a	false
sourceColumn	String	Exact match identifier for source column name.	non-empty String	n/a
targetColumn	String	Preferred field name to be used in the target.	Not Empty String	n/a

Property name	Type	Description	Constraints	Default
type	String	<p>Identifies the source column's data type which is used to transform the column value to a properly formatted String within the target. Supported types include: [STRING, VARCHAR, TEXT, INT, DATE, DATE_TIME, BYTE, DEC, NUMERIC, DECIMAL, DOUBLE, FLOAT, LONG, SHORT, RAW, BLOB, CLOB, HASHMAP, CUSTOM]</p> <p>CUSTOM data type is unique in that it bypasses column value transformation to a String which allows it to be converted manually within a Custom Stage. An example would be converting to a proprietary Oracle Timestamp format. Failure to convert this data type manually will cause errors in Redis-based sinks.</p>	<p>regexp = " [\\s<>(){}\\[\\]\\\"'\\/\\\\\\\\;` !@#\$\$%&*]*\$"</p>	STROING

Property name	Type	Description	Constraints	Default
changedColumnOnlyEnabled	Boolean	When enabled, only allows changed (delta) column values to be replicated to the target unless targetKey is enabled. When changedColumnsOnlyEnabled=true at the table-level, this flag will be overridden. This is currently only supported for RDB sources.	n/a	false
passThroughEnabled	Boolean	When disabled, the source column value will not be published to the pipeline therefore it cannot be accessed within a custom stage nor any sink. The purpose of this flag is to allow source column values to be used for targetKey composition without adding the column's name/value pair as a field within the target. As an example, this is common for sources like MongoDB which generate a "_id" key which can be used as a targetKey but has no value as a field.	n/a	true

Property name	Type	Description	Constraints	Default
index	Integer	This is currently for metadata purposes only and has no functional value.	n/a	n/a
dateFormat	String	Used by DATE and DATE_TIME type to override their default. Default formats are as follows: DATE = YYYY-MM-dd DATE_TIME = YYYY-MM-dd HH:mm:ss.S		n/a
nullFormat	String	Users can define how a column value=NULL will be represented in the target.	n/a	Default is an EMPTY String.

Job Source Table Initial Load Properties

Table 14. Job Source Table Column Properties

Property name	Type	Description	Constraints	Default
partitions	Integer	<p>Indicates how many partitions to create during startJob process.</p> <p>This attribute is ONLY used to partition an initial load with jobType=LOAD.</p> <p>Each table should be partitioned based on its own size and release window SLAs.</p> <p>It's common practice to leverage more partitions for an initial load than on streaming. Please see the Production Readiness section for more detail. Disclaimer: If the source table has fewer than 500 rows, which is common in a POC/dev environment, all but partition:1 will be stopped so all the rows are loaded from a single partition.</p>	min=1	1

Property name	Type	Description	Constraints	Default
maxPartitionsPerClusterMember	Integer	<p>Limits how many task partitions can be claimed, and executed, multi-tenant on the same Redis Connect instance (JVM).</p> <p>If the limit forces partitions to span more nodes than are currently deployed, then the initial load will queue the instantiation of tasks until capacity is reallocated (e.g. earlier tasks complete their load partition).</p> <p>This is not a job-level limit; it is only specific at the table level. 0 represents no limit.</p>	min=0	0

Property name	Type	Description	Constraints	Default
customWhereClause	String	<p>Users can specify a WHERE clause to filter the rows required for initial load. Only the following sources are supported:</p> <ul style="list-style-type: none"> - RDB sources support JDBC compliant WHERE statements - MongoDB supports a BSON filter - Gemfire supports an Apache Geode WHERE Clause 		
rowIndexUsedAsTargetKey Enabled	Boolean	RDB sources can have tables without primary keys. For those cases, rowIndex can be used as a unique identifier for partitioning purposes. This is only supported for RDB sources and only for initial load only / ETL jobs.		false

Database Properties

Table 15. Database Properties

Property name	Type	Description	Default
connectionType	String	<p>Distinguishes between Job Manager, Job Source, Job Target, and Job Checkpoint databases.</p> <p>This field is auto-generated.</p>	
databaseType	Enum	<p>The following database types are supported:</p> <p>[DB2, FILES, GEMFIRE, MONGODB, PostgreSQL, ORACLE, POSTGRES, REDIS, REDIS_STREAMS_MESSAGE_BROKER, SPLUNK, SQL_SERVER, VERTICA]</p> <p>NONE is used for custom stages. Also see userDefinedType.</p> <p>This is a required field.</p>	
databaseURL	String	<p>Database URLs specify the subprotocol (the database connectivity mechanism), the database, or server identifier, and a list of properties.</p> <p>Not required for Gemfire, Splunk, and Files databaseTypes.</p>	

Property name	Type	Description	Default
credentialsDirectoryPath		<p>Each database type other than NONE (used for custom jobs) requires a credentials property file, even if the credentials are not required. Credentials property files must adhere to the following filename pattern: redisconnect_credentials_{source_type target_type}_{job_name}.properties</p> <p>The only exception is for Job Manager which has fixed name: redisconnect_credentials_jobmanager.properties</p>	

Property name	Type	Description	Default
credentialsRotationEventListenerEnabled	Boolean	<p>When enabled, the credentialsDirectoryPath will be periodically scanned for changes that are specific to the property file associated with this database.</p> <p>If a change is identified, the listener will create a new connection without bringing down the Redis Connect instance nor stopping the job. There might be a momentary pause in pipeline processing while the connection is being reestablished. No data will be lost in this process.</p> <p>Disclaimer: if the new credentials cannot be used to create a connection, the job will be stopped for manual intervention.</p>	false
credentialsRotationEventListenerInterval	Integer	Fixed rate scheduled interval that scans the credentialsDirectoryPath for changes when credentialsRotationEventListenerEnabled is enabled. Measured in milliseconds. Minimum=60000	60000
customConfiguration	See Custom Configurations.	Map <String, String>	